

TASK 1:

Storage – HDD/SSD
Ex: Corsair Neutron XT 1.92 TB

Aftermarket cooling setup
Ex: Corsair AF120 ELITE

Cabinet
Ex: Corsair 4000D Airflow
ATX

GPU
Ex: MSI GeForce
RTX 3090 TI

RAM (2x)
Ex: Kingston FURY
Beast DDR5-5600

PSU/SMPS
Ex: EVGA SuperNOVA 1600
P2 1600 W

Motherboard
Ex: MSI MAG B650M MORTAR WIFI Micro ATX

CPU (+ CPU Fan)
Ex: AMD Ryzen 9 7950X

Image Source: <https://www.techpowerup.com/img/17-01-06/7bb8a0d2977d.jpg>

Task 2:

Many Deep Learning based models are “CUDA-based research”. One of the more recent CUDA-based research that I find interesting is the “Attention Is All You Need[1]” paper.

Why I find it interesting:

- This paper introduced the “Transformer” architecture, which is now one of the primary model architecture of choice.
- This architecture allows massive parallelisation – really putting to use the numerous CUDA cores of GPUs.
- This architecture makes it really possible to do NLP, Image and multiple domains in a unified architecture.

[1] Vaswani, A. (2017). Attention is all you need. Advances in Neural Information Processing Systems.

Name: Ramadoss, Adhithyan;
Matriculation Number: 22322010
IdM: um41adoh